

ML_Lab2

March 14, 2020

0.0.1 1. Find the url for the California Test Score Data Set from the following website:

<https://vincentarelbundock.github.io/Rdatasets/datasets.html>

Read through the “DOC” file to understand the variables in the dataset, then use the following url to import the data.

<https://vincentarelbundock.github.io/Rdatasets/csv/Ecdat/Caschool.csv>

The target data (i.e. the dependent variable) is named “testscr”. You can use all variables in the data except for “readscr” and “mathscr” in the following analysis. (These two variables were used to generate the dependent variable).

```
[79]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

import sys
import warnings

if not sys.warnoptions:
    warnings.simplefilter("ignore")
```

```
[2]: data = pd.read_csv("https://vincentarelbundock.github.io/Rdatasets/csv/Ecdat/
→Caschool.csv")
```

0.0.2 1.1 Visualize the univariate distribution of the target feature and each of the three continuous explanatory variables that you think are likely to have a relationship with the target feature.

```
[3]: data.head()
```

```
[3]: Unnamed: 0  distcod  county  district grspan \
0          1    75119  Alameda  Sunol Glen Unified  KK-08
1          2    61499   Butte  Manzanita Elementary  KK-08
2          3    61549   Butte  Thermalito Union Elementary  KK-08
3          4    61457   Butte  Golden Feather Union Elementary  KK-08
4          5    61523   Butte  Palermo Union Elementary  KK-08

enrltot  teachers  calwpct  mealpct  computer  testscr  compstu \
0      195  10.900000  0.510200  2.040800      67  690.799988  0.343590
```

1	240	11.150000	15.416700	47.916698	101	661.200012	0.420833
2	1550	82.900002	55.032299	76.322601	169	643.599976	0.109032
3	243	14.000000	36.475399	77.049202	85	647.700012	0.349794
4	1335	71.500000	33.108601	78.427002	171	640.849976	0.128090

	expnstu	str	avginc	elpct	readscr	mathscr
0	6384.911133	17.889910	22.690001	0.000000	691.599976	690.000000
1	5099.380859	21.524664	9.824000	4.583333	660.500000	661.900024
2	5501.954590	18.697226	8.978000	30.000002	636.299988	650.900024
3	7101.831055	17.357143	8.978000	0.000000	651.900024	643.500000
4	5235.987793	18.671329	9.080333	13.857677	641.799988	639.900024

```
[4]: ca_schools = data.drop(['readscr', 'mathscr', 'distcod'], axis=1)
ca_schools.head()
```


```
[4]: Unnamed: 0    county                                district grspan  enrltot  \
0           1  Alameda                        Sunol Glen Unified  KK-08      195
1           2    Butte                Manzanita Elementary  KK-08      240
2           3    Butte          Thermalito Union Elementary  KK-08     1550
3           4    Butte  Golden Feather Union Elementary  KK-08      243
4           5    Butte          Palermo Union Elementary  KK-08     1335
```

	teachers	calwpct	mealpct	computer	testscr	compstu	\
0	10.900000	0.510200	2.040800	67	690.799988	0.343590	
1	11.150000	15.416700	47.916698	101	661.200012	0.420833	
2	82.900002	55.032299	76.322601	169	643.599976	0.109032	
3	14.000000	36.475399	77.049202	85	647.700012	0.349794	
4	71.500000	33.108601	78.427002	171	640.849976	0.128090	

	expnstu	str	avginc	elpct
0	6384.911133	17.889910	22.690001	0.000000
1	5099.380859	21.524664	9.824000	4.583333
2	5501.954590	18.697226	8.978000	30.000002
3	7101.831055	17.357143	8.978000	0.000000
4	5235.987793	18.671329	9.080333	13.857677


```
[5]: #Univariate Distribution: Test Scores
df_test = pd.DataFrame(ca_schools, columns=["testscr"])
df_test
sns.distplot(df_test)
```

```
[5]: <matplotlib.axes._subplots.AxesSubplot at 0x1d93e7374e0>
```




```
[6]: #Univariate Distribution: Computers Per Student
df_comp = pd.DataFrame(ca_schools, columns=["compstu"])
df_comp
sns.distplot(df_comp)
```

```
[6]: <matplotlib.axes._subplots.AxesSubplot at 0x1d93ea7af28>
```




```
[7]: #Univariate Distribution: Student/Teacher Ratio
df_str = pd.DataFrame(ca_schools, columns=["str"])
df_str
sns.distplot(df_str)
```

[7]: <matplotlib.axes._subplots.AxesSubplot at 0x1d93eb34390>



```
[8]: #Univariate Distribution: Average Income
df_inc = pd.DataFrame(ca_schools, columns=["avginc"])
df_inc
sns.distplot(df_inc)
```


[8]: <matplotlib.axes._subplots.AxesSubplot at 0x1d93eb3cc50>



0.03 1.2 Visualize the dependency of the target on each feature from 1.1.


```
[9]: sns.jointplot(x="compstu", y="testscr", data=ca_schools)
```

```
[9]: <seaborn.axisgrid.JointGrid at 0x1d93ec5b9e8>
```




```
[10]: sns.jointplot(x="str", y="testscr", data=ca_schools)
```

```
[10]: <seaborn.axisgrid.JointGrid at 0x1d93edc99e8>
```



```
[11]: sns.jointplot(x="avginc", y="testscr", data=ca_schools)
```

```
[11]: <seaborn.axisgrid.JointGrid at 0x1d93ef228d0>
```



0.0.4 1.3 Split data in training and test set. Build models that evaluate the relationship between all available X variables in the California test dataset and the target variable. Evaluate KNN for regression, Linear Regression (OLS), Ridge, and Lasso using cross-validation with the default parameters. Does scaling the data with the StandardScaler help?

Split data in training and test set

```
[32]: y = ca_schools['testscr']
      y=y.astype("int64")

[33]: X = ca_schools.drop(['Unnamed: 0', 'county', 'district','grspan','testscr'],
      ↪axis=1)
```



```
[34]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y)
```

```
[35]: from sklearn import preprocessing
scaler = preprocessing.StandardScaler().fit(X_train)
X_train_scaled = scaler.transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
[16]: from sklearn.model_selection import cross_val_score
from sklearn.model_selection import RepeatedKFold
from sklearn.model_selection import StratifiedKFold
from sklearn.model_selection import KFold

kfold = KFold(n_splits=5)
skfold = StratifiedKFold(n_splits=5, shuffle=True)
rkf = RepeatedKFold(n_splits=5, n_repeats=10)
```

0.0.5 KNN for Regression

```
[17]: from sklearn.neighbors import KNeighborsRegressor

knn = KNeighborsRegressor(n_neighbors=5)
knn.fit(X_train, y_train)

print("R^2: {:.2f}".format(knn.score(X_test, y_test)))

y_pred = knn.predict(X_test)
```

R²: 0.04

KNN for Regression w/ Scaled Data

```
[18]: knn = KNeighborsRegressor(n_neighbors=5)
knn.fit(X_train_scaled, y_train)

print("R^2: {:.2f}".format(knn.score(X_test_scaled, y_test)))
```

R²: 0.69

Scaling the data using the Standard Scalar, which standardizes by calculating a standard score. This score is calculated through the use of the following formula: $z = (x-u)/s$. Where z is the score, x is the unit to be standardized u is the mean of the data samples and s is their standard deviation. In the case of KNN Regression standardizing the data has produced a model with greater explanatory power, as we can see an increase in R^2 from 0.04 to 0.69 suggesting that when standardized 69% of the variation in y can be explained by x , where only 4% can be explained given no standardization.

Sources: 1. <https://scikit-learn.org/stable/modules/generated/sklearn.preprocessing.StandardScaler.html>
2. QMSSGR5073 'Preprocessing Data.ipynb', "Using StandardScaler()

KNN for Regression Cross Validation

```
[44]: #Empty space in regressor assumes n_neighbors = 5

print("KFold Mean:\n{}".format(
    cross_val_score(KNeighborsRegressor(), X, y, cv=kfold).mean()))

print("StratifiedKFold Mean:\n{}".format(
    cross_val_score(KNeighborsRegressor(), X, y, cv=skfold).mean()))

print("RepeatedKFold Mean:\n{}".format(
    cross_val_score(KNeighborsRegressor(), X, y, cv=rkf).mean()))
```

```
KFold Mean:
-16.071366760448093
StratifiedKFold Mean:
0.01232438589727749
RepeatedKFold Mean:
-0.017588968334293246
```

We see that when apply different types of cross-validation, holding the `n_neighbors` constant, we get back some negative scores (these scores equate to the average R^2 for the runs of cross validation) Upon running the model with the test and train data the negatives went away, however no examples on scikit or the notes suggest using the split data. Found examples use the data as a whole separated by `X` (independent variables) and `y` (dependent variables). As a result, the previous cell was provided, despite the negative measure of R^2 . Do note the stratified KFold mean is positive, which appears to suggest that we can get some small positive explanatory power by ensuring that each fold has the same proportion of observations.

Sources: 1. https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.cross_val_score.html
2. QMSSGR5073 "Knn_CV_and_Model_Tuning.ipynb, 'Using Cross validation for model evaluation' "

3. <https://towardsdatascience.com/cross-validation-explained-evaluating-estimator-performance-e51e5430ff85>

0.0.6 Linear Regression (OLS)

```
[45]: from sklearn.linear_model import LinearRegression
import statsmodels.api as sm
```

```
[46]: X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42)
```

Linear Regression (OLS) w/ Non-Scaled Data

```
[47]: lr = LinearRegression().fit(X_train, y_train)
```

```
[48]: X_train_new = sm.add_constant(X_train)
model = sm.OLS(y_train, X_train_new ).fit()

model.summary()
```

```
[48]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                OLS Regression Results
=====
Dep. Variable:          testscr      R-squared:                0.805
Model:                  OLS          Adj. R-squared:           0.799
Method:                 Least Squares   F-statistic:              125.5
Date:                  Mon, 02 Mar 2020   Prob (F-statistic):       1.24e-101
Time:                  11:32:12          Log-Likelihood:           -1099.0
No. Observations:      315              AIC:                     2220.
Df Residuals:          304              BIC:                     2261.
Df Model:              10
Covariance Type:       nonrobust
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
const         663.1277     10.936     60.635     0.000     641.607     684.648
enrltot         0.0002      0.002      0.095     0.924      -0.003      0.004
teachers       -0.0080      0.038     -0.208     0.835      -0.083      0.067
calwpct       -0.1048      0.068     -1.542     0.124      -0.239      0.029
mealpct       -0.3639      0.042     -8.745     0.000      -0.446     -0.282
computer        0.0022      0.003      0.661     0.509      -0.004      0.009
compstu       -1.3673      8.528     -0.160     0.873     -18.149     15.415
expnstu         0.0019      0.001      1.911     0.057     -5.65e-05      0.004
str           -0.3460      0.367     -0.942     0.347      -1.069      0.377
avginc         0.5201      0.096      5.395     0.000      0.330      0.710
elpct         -0.1963      0.042     -4.681     0.000      -0.279     -0.114
=====
Omnibus:                 1.564   Durbin-Watson:           1.873
Prob(Omnibus):           0.457   Jarque-Bera (JB):         1.377
Skew:                    0.016   Prob(JB):                 0.502
Kurtosis:                3.322   Cond. No.                 1.58e+05
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.58e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
"""
```

Linear Regression w/ Scaled Data

```
[49]: X_train_new = sm.add_constant(X_train_scaled)
      model = sm.OLS(y_train, X_train_new).fit()

      model.summary()
```

```
[49]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                OLS Regression Results
=====
Dep. Variable:                  testscr      R-squared:                  0.009
Model:                            OLS      Adj. R-squared:              -0.023
Method:                    Least Squares      F-statistic:                  0.2893
Date:                Mon, 02 Mar 2020      Prob (F-statistic):          0.983
Time:                      11:32:27      Log-Likelihood:             -1354.9
No. Observations:                315      AIC:                        2732.
Df Residuals:                    304      BIC:                        2773.
Df Model:                          10
Covariance Type:                nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	653.8063	1.024	638.376	0.000	651.791	655.822
x1	5.5483	15.980	0.347	0.729	-25.898	36.995
x2	-5.8866	17.138	-0.343	0.731	-39.610	27.837
x3	-0.1447	1.728	-0.084	0.933	-3.544	3.255
x4	-1.4030	2.539	-0.552	0.581	-6.400	3.594
x5	0.6326	3.486	0.181	0.856	-6.226	7.492
x6	-0.1151	1.241	-0.093	0.926	-2.557	2.326
x7	0.7187	1.429	0.503	0.615	-2.093	3.531
x8	-0.1529	1.539	-0.099	0.921	-3.182	2.876
x9	-1.3574	1.577	-0.861	0.390	-4.460	1.745
x10	1.8683	1.669	1.119	0.264	-1.417	5.153

```

=====
Omnibus:                        0.609      Durbin-Watson:              1.892
Prob(Omnibus):                  0.737      Jarque-Bera (JB):           0.713
Skew:                          0.011      Prob(JB):                   0.700
Kurtosis:                      2.768      Cond. No.                   43.4
=====

```

```
Warnings:
```

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly
specified.
"""
```

In the case of Linear Regression, scaling the data has a negative impact on the explanatory power of the model. We see some changes in the size of the coefficients, and the decrease of t-stats, which suggest the variables should not be included in the model, as one cannot reject the null hypothesis that the independent variables explain nothing with regard to our y (dependent variable).

Linear Regression Cross Validation

```
[53]: print("KFold:\n{}".format(
      cross_val_score(LinearRegression(), X, y, cv=kfold).mean()))

      print("StratifiedKFold:\n{}".format(
      cross_val_score(LinearRegression(), X, y, cv=skfold).mean()))

      print("RepeatedKFold:\n{}".format(
      cross_val_score(LinearRegression(), X, y, cv=rkf).mean()))
```

```
KFold:
-4.379031971136683
StratifiedKFold:
0.80008382405661
RepeatedKFold:
0.7926505547400525
```

We see that when apply different types of cross-validation we get back a negative score from KFold. Cross Validation through Stratified Kfold and Repeated Kfold suggest that the independent variables in the linear regression can explain 80% of the variation in our dependent variable.

0.0.7 Ridge Regression

```
[54]: from sklearn.linear_model import Ridge
      from sklearn.linear_model import RidgeCV
```

Ridge Regression w/ Non-Scaled Data

```
[65]: ridge = Ridge().fit(X_train, y_train)
      print("Training R^2: {:.2f}".format(ridge.score(X_train, y_train)))
      print("Test R^2: {:.2f}".format(ridge.score(X_test, y_test)))
```

```
Training R^2: 0.80
Test R^2: 0.81
```

Ridge Regression w/ Scaled Data

```
[62]: scaler = preprocessing.StandardScaler()
      scaler.fit(X_train)
      X_train_scaled = scaler.transform(X_train)
      X_test_scaled = scaler.transform(X_test)

[64]: ridge = Ridge().fit(X_train_scaled, y_train)
      print("Training R^2: {:.2f}".format(ridge.score(X_train_scaled, y_train)))
      print("Test R^2: {:.2f}".format(ridge.score(X_test_scaled, y_test)))
```

```
Training R^2: 0.80
Test R^2: 0.81
```

This ridge regression does not seem to experience a change in explanatory power given the scaled data.

Ridge Regression Cross Validation

```
[66]: print("KFold:\n{}".format(
      cross_val_score(Ridge(), X, y, cv=kfold).mean()))

      print("StratifiedKFold:\n{}".format(
      cross_val_score(Ridge(), X, y, cv=skfold).mean()))

      print("RepeatedKFold:\n{}".format(
      cross_val_score(Ridge(), X, y, cv=rkf).mean()))
```

```
KFold:
-4.3006576536144925
StratifiedKFold:
0.8047688959618053
RepeatedKFold:
0.7917275577377194
```

We see that when apply different types of cross-validation we get back a negative score from KFold. Cross Validation through Stratified Kfold and Repeated Kfold suggest that the independent variables in the ridge regression can explain roughly 80% of the variation in our dependent variable.

0.0.8 Lasso Regression

Lasso Regression w/ Non-Scaled Data

```
[68]: from sklearn.linear_model import Lasso

[69]: lasso = Lasso().fit(X_train, y_train)
      print("Training R^2: {:.2f}".format(lasso.score(X_train, y_train)))
      print("Test R^2: {:.2f}".format(lasso.score(X_test, y_test)))
      print("Number of features used: {}".format(np.sum(lasso.coef_ != 0)))
      #print("lasso.coef_: {}".format(lasso.coef_))
```

```
Training R^2: 0.80
Test R^2: 0.81
Number of features used: 8
```

Lasso Regression w/ Scaled Data

```
[71]: lasso = Lasso().fit(X_train_scaled, y_train)
      print("Training R^2: {:.2f}".format(lasso.score(X_train_scaled, y_train)))
      print("Test R^2: {:.2f}".format(lasso.score(X_test_scaled, y_test)))
      print("Number of features used: {}".format(np.sum(lasso.coef_ != 0)))

      #print("lasso.coef_: {}".format(lasso.coef_))
```

```
Training R^2: 0.80
Test R^2: 0.79
Number of features used: 5
```

This lasso regression seems to experience a change in explanatory power given the scaled data, but it is very small and negative. This change is only with regard to the explanatory power on the testing data.

Lasso Regression Cross Validation

```
[72]: print("KFold:\n{}".format(
      cross_val_score(Lasso(), X, y, cv=kfold).mean()))

      print("StratifiedKFold:\n{}".format(
      cross_val_score(Lasso(), X, y, cv=skfold).mean()))

      print("RepeatedKFold:\n{}".format(
      cross_val_score(Lasso(), X, y, cv=rkf).mean()))
```

```
KFold:
-4.28885690538308
StratifiedKFold:
0.8009257618920553
RepeatedKFold:
0.7937523019896281
```

We see that when apply different types of cross-validation we get back a negative score from KFold. Cross Validation through Stratified Kfold and Repeated Kfold suggest that the independent variables in the lasso regression can explain roughly 80% of the variation in our dependent variable.

0.0.9 1.4 Tune the parameters of the models where possible using GridSearchCV. Do the results improve?

```
[76]: from sklearn.pipeline import make_pipeline
      from sklearn.model_selection import GridSearchCV
      from sklearn.preprocessing import StandardScaler
```

KNN GridSearch CV

```
[77]: knn_pipe = make_pipeline(KNeighborsRegressor())

[80]: param_grid = {'kneighborsregressor__n_neighbors': range(1, 10)}
      grid = GridSearchCV(knn_pipe, param_grid, cv=10)
      grid.fit(X_train, y_train)
      print(grid.best_params_)
      print(grid.best_score_)
```

```
{'kneighborsregressor__n_neighbors': 9}
-0.016257788001062312
```

```
[81]: knn_pipe = make_pipeline(StandardScaler(), KNeighborsRegressor())
```

```
[82]: param_grid = {'kneighborsregressor__n_neighbors': range(1, 10)}
grid = GridSearchCV(knn_pipe, param_grid, cv=10)
grid.fit(X_train, y_train)
print(grid.best_params_)
print(grid.best_score_)
```

```
{'kneighborsregressor__n_neighbors': 7}
0.7172474999397614
```

GridSearch CV appears to suggest that for the KNN regressor benefits from scaling, as the R^2 has become positive. the ideal `n_neighbors` would be 7 and this would produce a model with an explanatory power of 0.717, which suggests that 72% of the variation in y (dependent variable) can be explained by x (independent variables)

Ridge Regression GridSearch CV

```
[83]: ridge_pipe = make_pipeline(Ridge())
```

```
[84]: param_grid = {'ridge__alpha': [0.01, 0.1, 1, 10, 100,1000,10000]}
grid = GridSearchCV(ridge_pipe, param_grid, cv=10)
grid.fit(X_train, y_train)
print(grid.best_params_)
print(grid.best_score_)
```

```
{'ridge__alpha': 1000}
0.7822304800155856
```

```
[85]: ridge_pipe = make_pipeline(StandardScaler(), Ridge())
```

```
[86]: param_grid = {'ridge__alpha': [0.01, 0.1, 1, 10, 100,1000,10000]}
grid = GridSearchCV(ridge_pipe, param_grid, cv=10)
grid.fit(X_train, y_train)
print(grid.best_params_)
print(grid.best_score_)
```

```
{'ridge__alpha': 10}
0.7828517182362517
```

Our grid search seems to suggest that the best model will have an alpha of 1000 with unscaled data and 10 with scaled data. This appears to make sense because as the model's alpha approaches 0 the model begins to become increasingly similar to an OLS regression. So for data that is unscaled a larger alpha makes sense as it requires more of a smoothness constraint to find explanatory power. The scaled data can operate more like an OLS because the variation in the data has been standardized, thus making it simpler for the OLS to predict.

```
[92]: ridge10 = Ridge(alpha=10).fit(X_train, y_train)
print("Training  $R^2$ : {:.2f}".format(ridge10.score(X_train, y_train)))
print("Test set  $R^2$ : {:.2f}".format(ridge10.score(X_test, y_test)))
```



```
Training R^2: 0.80
Test set R^2: 0.81
```

```
[91]: ridge1000 = Ridge(alpha=1000).fit(X_train_scaled, y_train)
print("Training R^2: {:.2f}".format(ridge1000.score(X_train_scaled, y_train)))
print("Test set R^2: {:.2f}".format(ridge1000.score(X_test_scaled, y_test)))
```

```
Training R^2: 0.56
Test set R^2: 0.52
```

Lasso Regression GridSearch CV

```
[93]: lasso_pipe = make_pipeline(Lasso())
```

```
[94]: param_grid = {'lasso__alpha': [0.01, 0.1, 1, 10, 100, 1000, 10000]}
grid = GridSearchCV(lasso_pipe, param_grid, cv=10)
grid.fit(X_train, y_train)
print(grid.best_params_)
print(grid.best_score_)
```

```
{'lasso__alpha': 1}
0.7823411438316094
```

```
[95]: lasso_pipe = make_pipeline(StandardScaler(), Lasso())
```

```
[96]: param_grid = {'lasso__alpha': [0.01, 0.1, 1, 10, 100, 1000, 10000]}
grid = GridSearchCV(lasso_pipe, param_grid, cv=10)
grid.fit(X_train, y_train)
print(grid.best_params_)
print(grid.score(X_test, y_test))
```

```
{'lasso__alpha': 0.1}
0.806544494469053
```

With regard to lasso, our grid search seems to suggest that the best model will have an alpha of 1 with unscaled data and 0.1 with scaled data. This appears to make sense because as the model's alpha approaches 0 the model begins to become increasingly similar to an OLS regression, but as alpha becomes too large coefficients will begin to 0 out. It is not surprising then that we see much lower alpha measures for lasso than for ridge while keeping a similar level of explanatory power.

```
[98]: lasso01 = Lasso(alpha=0.1).fit(X_train, y_train)
print("Training R^2: {:.2f}".format(lasso01.score(X_train, y_train)))
print("Test R^2: {:.2f}".format(lasso01.score(X_test, y_test)))
```

```
Training R^2: 0.80
Test R^2: 0.81
```

```
[105]: int(lasso.score(X_test, y_test))
print(lasso.coef_)
```

```
[ -0.          -0.          -0.          -10.6393318  -0.
   0.          0.46691845 -0.19918508   3.34742122 -2.53863698]
```

One can see from the basic lasso coefficients that the previously mentioned zeroing out has occurred.

0.0.10 1.5 Compare the coefficients of your two best linear models (not knn), do they agree on which features are important?

```
[103]: print(ridge10.score(X_test, y_test))
print(ridge10.coef_)
```

```
0.8088485493057405
[ 1.18796341e-04 -6.54095365e-03 -1.03943211e-01 -3.64516588e-01
  1.93892193e-03 -1.09733165e-01  1.88222319e-03 -3.30166182e-01
  5.18872011e-01 -1.95405911e-01]
```

```
[104]: print(lr.score(X_test, y_test))
print(lr.coef_)
```

```
0.8081674788950314
[ 1.64271449e-04 -7.97874165e-03 -1.04839109e-01 -3.63865590e-01
  2.15309332e-03 -1.36726004e+00  1.88904899e-03 -3.45970527e-01
  5.20091618e-01 -1.96277441e-01]
```

The models provided above have similar explanatory power, suggesting that their best models will provide an R^2 value near .80 which suggests that 80% of the variation in our dependent variable can be explained by our independent variables. We also can see that while the magnitude of the coefficients has changed slight from one model to the next, the signs stay the same, this is good as the models don't contradict the explanatory nature of the features within them. Realizing similar scores across different types of regressions is promising, as it suggests that the data does have the ability to explain the dependent variables variation in a variety of settings.

0.0.11 1.6 Discuss which final model you would choose to predict new data

I would choose the ridge regression having gone through the analysis so far. The ridge regression leverages the L2 penalty which makes it more challenging to discover a signal by adding "a value equal to the square of the magnitude of coefficients" to the model. Ridge also allows a parsimonious model to be created despite the presence of multicollinearity, while this HW does not discuss multicollinearity, if I were doing research I would certainly consider a model that had strong explanatory power while also having mitigated the potential risk of finding multicollinearity.

Sources: 1. <https://www.statisticshowto.datasciencecentral.com/ridge-regression/> 2. <https://www.statisticshowto.datasciencecentral.com/regularization/>

0.0.12 Question 2

First, import the red and the white wine csv files into separate pandas dataframes from the following website:

<https://archive.ics.uci.edu/ml/machine-learning-databases/wine-quality>

(Note: you need to adjust the argument for `read_csv()` from `sep=','` to `sep=';'`)

Add a new column to each data frame called "winetype". For the white wine dataset label the values in this column with a 0, indicating white wine. For the red wine dataset, label values with a 1, indicating red wine. Combine both datasets into a single dataframe.

The target data (i.e. the dependent variable) is "winetype".

```
[3]: import pandas as pd
```

```
[106]: red_wine = pd.read_csv("D:/QMSS/Spring/Machine Learning/winequality-red.  
→csv", sep=';')  
white_wine = pd.read_csv("D:/QMSS/Spring/Machine Learning/winequality-white.  
→csv", sep=';')
```

```
[107]: red_wine['winetype'] = 1  
white_wine['winetype'] = 0  
frames = [red_wine, white_wine]  
wines = pd.concat(frames)  
wines.head()
```

```
[107]:  fixed acidity  volatile acidity  citric acid  residual sugar  chlorides  \  
0           7.4             0.70         0.00           1.9        0.076  
1           7.8             0.88         0.00           2.6        0.098  
2           7.8             0.76         0.04           2.3        0.092  
3          11.2             0.28         0.56           1.9        0.075  
4           7.4             0.70         0.00           1.9        0.076
```


```
    free sulfur dioxide  total sulfur dioxide  density  pH  sulphates  \  
0              11.0              34.0  0.9978  3.51         0.56  
1              25.0              67.0  0.9968  3.20         0.68  
2              15.0              54.0  0.9970  3.26         0.65  
3              17.0              60.0  0.9980  3.16         0.58  
4              11.0              34.0  0.9978  3.51         0.56
```

```
    alcohol  quality  winetype  
0       9.4        5         1  
1       9.8        5         1  
2       9.8        5         1  
3       9.8        6         1  
4       9.4        5         1
```

0.0.13 2.1 Visualize the univariate distribution of the target feature and each of the three explanatory variables that you think are likely to have a relationship with the target feature.


```
[108]: #Univariate Distribution Wine Types
wine_type = pd.DataFrame(wines, columns=["winetype"])
sns.distplot(wine_type)
```

```
[108]: <matplotlib.axes._subplots.AxesSubplot at 0x1d9427479e8>
```




```
[109]: #Univariate Distribution: Alcohol Content by Percentage
alcohol = pd.DataFrame(wines, columns=["alcohol"])
sns.distplot(alcohol)
```

```
[109]: <matplotlib.axes._subplots.AxesSubplot at 0x1d93f9cc1d0>
```




```
[110]: #Univariate Distribution: Sugar Content
sugar = pd.DataFrame(wines, columns=["residual sugar"])
sns.distplot(sugar)
```

```
[110]: <matplotlib.axes._subplots.AxesSubplot at 0x1d942897f98>
```



```
[111]: #Univariate Distribution: Quality Rating
quality = pd.DataFrame(wines, columns=["quality"])
sns.distplot(quality)
```

```
[111]: <matplotlib.axes._subplots.AxesSubplot at 0x1d94298ad30>
```



0.0.14 2.2 Split data into training and test set. Build models that evaluate the relationship between all available X variables in the dataset and the target variable. Evaluate Logistic Regression, Penalized Logistic Regression, and KNN for classification using cross-validation. How different are the results? How does scaling the data with StandardScaler influence the results?

```
[112]: y = wines['winetype']
X = wines.loc[:, wines.columns!="winetype"]
```

```
[113]: X_train, X_test, y_train, y_test = train_test_split(X, y)
```

```
[114]: #from sklearn import preprocessing
scaler = preprocessing.StandardScaler().fit(X_train)
X_train_scaled = scaler.transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

KNN for Classification w/ Non-Scaled Data

```
[115]: from sklearn.neighbors import KNeighborsClassifier

knn = KNeighborsClassifier(n_neighbors=5)
```

```
knn.fit(X_train, y_train)

print("accuracy: {:.2f}".format(knn.score(X_test, y_test)))

y_pred = knn.predict(X_test) # y_pred includes your predictions
```

accuracy: 0.94

KNN for Classification w/ Scaled Data

```
[116]: knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train_scaled, y_train)

print("accuracy: {:.2f}".format(knn.score(X_test_scaled, y_test)))

y_pred = knn.predict(X_test_scaled)
```

accuracy: 0.99

We see that for KNN Classification, scaling the data has made produced a model that has stronger predictive power, with classification accuracy increasing from 0.94 to 0.99.

Cross Validation for KNN Classification

```
[118]: print("KFold:\n{}".format(
cross_val_score(KNeighborsClassifier(), X, y, cv=kfold).mean()))

print("StratifiedKFold:\n{}".format(
cross_val_score(KNeighborsClassifier(n_neighbors=5), X, y, cv=skfold).mean()))

print("RepeatedKFold:\n{}".format(
cross_val_score(KNeighborsClassifier(n_neighbors=5), X, y, cv=rkf).mean()))
```

KFold:
0.9089049564753953
StratifiedKFold:
0.9435133419704602
RepeatedKFold:
0.942158429561201

We see that when cross validating, using a variety of different approaches, we are still receiving measure that show strong classification abilities with all means being greater than 0.9.

0.0.15 Logistic Regression

```
[119]: from sklearn.linear_model import LogisticRegression
```

Logistic Regression w/ Non-Scaled Data

```
[120]: logreg = LogisticRegression(C=1e90).fit(X_train, y_train)
```

```
[121]: import statsmodels.api as sm
```

```
X_train_new = sm.add_constant(X_train)
```

```
model = sm.GLM(y_train, X_train_new, family=sm.families.Binomial()).fit()
```

```
model.summary()
```

```
[121]: <class 'statsmodels.iolib.summary.Summary'>
```

```
"""
```

Generalized Linear Model Regression Results

```
=====
Dep. Variable:          winetype    No. Observations:          4872
Model:                  GLM         Df Residuals:              4859
Model Family:           Binomial    Df Model:                  12
Link Function:          logit       Scale:                    1.0000
Method:                 IRLS        Log-Likelihood:           -168.26
Date:                   Mon, 02 Mar 2020    Deviance:                 336.53
Time:                   13:03:59    Pearson chi2:             1.16e+07
No. Iterations:         10
Covariance Type:        nonrobust
=====
```

```
=====
                                coef    std err          z      P>|z|      [0.025
0.975]
-----
const                -1643.1565    205.302     -8.004     0.000    -2045.540
-1240.773
fixed acidity         -0.0882      0.264     -0.334     0.739     -0.606
0.430
volatile acidity       7.6527      1.254      6.103     0.000      5.195
10.110
citric acid           -2.0395      1.281     -1.592     0.111     -4.551
0.472
residual sugar        -0.9643      0.115     -8.416     0.000     -1.189
-0.740
chlorides             21.0194      4.453      4.720     0.000     12.292
29.747
free sulfur dioxide    0.0675      0.015      4.442     0.000      0.038
0.097
total sulfur dioxide   -0.0512      0.006     -9.183     0.000     -0.062
-0.040
density              1636.5619    209.430      7.814     0.000    1226.086
2047.038
pH                   -0.8112      1.597     -0.508     0.612     -3.942
=====
```



```

2.319
sulphates          3.1885      1.420      2.245      0.025      0.405
5.972
alcohol           1.5839      0.306      5.177      0.000      0.984
2.184
quality           0.3365      0.218      1.542      0.123     -0.091
0.764
=====
=====
"""

```

```

[129]: print("Training Pseudo R^2: {:.3f}".format(logreg.score(X_train, y_train)))
       print("Test Pseudo R^2: {:.3f}".format(logreg.score(X_test, y_test)))

```

```

Training Pseudo R^2: 0.989
Test Pseudo R^2: 0.988

```

Logistic Regression w/ Scaled Data

```

[125]: logregsc = LogisticRegression(C=1e90).fit(X_train_scaled, y_train)

[126]: X_train_new = sm.add_constant(X_train_scaled) #adding a column of 1s to the
        ↪matrix

model = sm.GLM(y_train, X_train_new, family=sm.families.Binomial()).fit()

model.summary()

```

```

[126]: <class 'statsmodels.iolib.summary.Summary'>
      """
          Generalized Linear Model Regression Results
      =====
      Dep. Variable:          winetype      No. Observations:          4872
      Model:                  GLM           Df Residuals:            4859
      Model Family:           Binomial       Df Model:                  12
      Link Function:          logit          Scale:                    1.0000
      Method:                 IRLS          Log-Likelihood:           -168.26
      Date:                   Mon, 02 Mar 2020      Deviance:                 336.53
      Time:                   13:07:06            Pearson chi2:             1.16e+07
      No. Iterations:         10
      Covariance Type:        nonrobust
      =====

```

	coef	std err	z	P> z	[0.025	0.975]
const	-4.2940	0.332	-12.921	0.000	-4.945	-3.643
x1	-0.1133	0.340	-0.334	0.739	-0.779	0.552
x2	1.2588	0.206	6.103	0.000	0.855	1.663
x3	-0.2957	0.186	-1.592	0.111	-0.660	0.068
x4	-4.6024	0.547	-8.416	0.000	-5.674	-3.531

x5	0.7126	0.151	4.720	0.000	0.417	1.009
x6	1.2129	0.273	4.442	0.000	0.678	1.748
x7	-2.8843	0.314	-9.183	0.000	-3.500	-2.269
x8	4.9373	0.632	7.814	0.000	3.699	6.176
x9	-0.1304	0.257	-0.508	0.612	-0.634	0.373
x10	0.4657	0.207	2.245	0.025	0.059	0.872
x11	1.8882	0.365	5.177	0.000	1.173	2.603
x12	0.2923	0.190	1.542	0.123	-0.079	0.664

=====
 ""

```
[128]: print("Training Pseudo R^2: {:.3f}".format(logregsc.score(X_train_scaled,
→y_train)))
print("Test Pseudo R^2: {:.3f}".format(logregsc.score(X_test_scaled, y_test)))
```

Training Pseudo R^2: 0.994

Test Pseudo R^2: 0.995

We can see that for logistic regression scaling the data has produced a model that has slightly increased explanatory power, with the independent variables being able to explain 99% of the variation in y. Our previous model could explain just under 99%.

Logistic Regression Cross Validation

```
[130]: print("KFold:\n{}".format(
cross_val_score(LogisticRegression(C=1e90), X, y, cv=kfold).mean()))

print("StratifiedKFold:\n{}".format(
cross_val_score(LogisticRegression(C=1e90), X, y, cv=skfold).mean()))

print("RepeatedKFold:\n{}".format(
cross_val_score(LogisticRegression(C=1e90), X, y, cv=rkf).mean()))
```

KFold:

0.981072777876473

StratifiedKFold:

0.9881478036322833

RepeatedKFold:

0.9883638064783561

We see that when cross validating, using a variety of different approaches, we are still receiving measure that show strong predicative abilities with all means being greater than 0.98.

0.0.16 Penalized Logistic Regression

Penalized Logistic Regression w/ Non-Scaled Data

```
[131]: lgr1 = LogisticRegression(C=1e90, penalty = 'l1').fit(X_train, y_train)
```

```
[132]: print("Training Pseudo R^2: {:.3f}".format(lgr1.score(X_train, y_train)))
print("Test Pseudo R^2: {:.3f}".format(lgr1.score(X_test, y_test)))
```

Training Pseudo R²: 0.989
Test Pseudo R²: 0.985

```
[150]: lgr2 = LogisticRegression(C=1e90, penalty = 'l2').fit(X_train, y_train)
```

```
[151]: print("Training Pseudo R^2: {:.3f}".format(lgr2.score(X_train, y_train)))
print("Test Pseudo R^2: {:.3f}".format(lgr2.score(X_test, y_test)))
```

Training Pseudo R²: 0.989
Test Pseudo R²: 0.988

Penalized Logistic Regression w/ Scaled Data

```
[152]: lgr1 = LogisticRegression(C=1e90, penalty = 'l1').fit(X_train_scaled, y_train)
```

```
[153]: print("Training Pseudo R^2: {:.3f}".format(lgr1.score(X_train_scaled, y_train)))
print("Test Pseudo R^2: {:.3f}".format(lgr1.score(X_test_scaled, y_test)))
```

Training Pseudo R²: 0.994
Test Pseudo R²: 0.995

```
[155]: lgr2 = LogisticRegression(C=1e90, penalty = 'l2').fit(X_train_scaled, y_train)
```

```
[156]: print("Training Pseudo R^2: {:.3f}".format(lgr2.score(X_train_scaled, y_train)))
print("Test Pseudo R^2: {:.3f}".format(lgr2.score(X_test_scaled, y_test)))
```

Training Pseudo R²: 0.994
Test Pseudo R²: 0.995

We can see that for penalized logistic regression scaling the data has produced a model that has slightly increased explanatory power, with the independent variables being able to explain 99% of the variation in y. Our previous model could explain just under 99%.

Penalized Logistic Regression Cross Validation (l1 Penalty)

```
[139]: print("KFold:\n{}".format(
cross_val_score(LogisticRegression(C=1e90,penalty="l1"), X, y, cv=kfold).
    ↳mean()))

print("StratifiedKFold:\n{}".format(
cross_val_score(LogisticRegression(C=1e90,penalty="l1"), X, y, cv=skfold).
    ↳mean()))

print("RepeatedKFold:\n{}".format(
cross_val_score(LogisticRegression(C=1e90,penalty="l1"), X, y, cv=rkf).mean()))
```

```
KFold:
0.9807650855687807
StratifiedKFold:
0.9870704062707037
RepeatedKFold:
0.9879790963463019
```

Penalized Logistic Regression Cross Validation (l2 Penalty)

```
[140]: print("KFold:\n{}".format(
cross_val_score(LogisticRegression(C=1e90,penalty="l2"), X, y, cv=kfold).
    ↳mean()))

print("StratifiedKFold:\n{}".format(
cross_val_score(LogisticRegression(C=1e90,penalty="l2"), X, y, cv=skfold).
    ↳mean()))

print("RepeatedKFold:\n{}".format(
cross_val_score(LogisticRegression(C=1e90,penalty="l2"), X, y, cv=rkf).mean()))
```

```
KFold:
0.981072777876473
StratifiedKFold:
0.9878404657150497
RepeatedKFold:
0.9882564576301297
```

We see that when accounting for penalty, cross validating, using a variety of different approaches, we are still receiving measure that show strong predicative abilities with all means being greater than 0.98.

0.0.17 2.3 Tune the parameters where possible using GridSearchCV. Do the results improve?

Logistic Regression w/ GridSearch CV

```
[178]: lgr = LogisticRegression()

penalty = ['l1', 'l2']

C = np.logspace(0, 4, 10)

hyperparameters = dict(C=C, penalty=penalty)
```

Logistic Regression w/ GridSearch CV Non-Scaled Data

```
[142]: #grid = GridSearchCV(lgr, param_grid, cv=10, scoring = 'accuracy', )
grid = GridSearchCV(lgr, param_grid=hyperparameters, cv=10)
```

```
grid_sc = GridSearchCV(lgr, param_grid=hyperparameters, cv=10)
```

```
[ ]:
```

```
[143]: grid.fit(X_train, y_train)
```

```
[143]: GridSearchCV(cv=10, error_score='raise-deprecating',
                  estimator=LogisticRegression(C=1.0, class_weight=None, dual=False,
                                                fit_intercept=True,
                                                intercept_scaling=1, l1_ratio=None,
                                                max_iter=100, multi_class='warn',
                                                n_jobs=None, penalty='l2',
                                                random_state=None, solver='warn',
                                                tol=0.0001, verbose=0,
                                                warm_start=False),
                  iid='warn', n_jobs=None,
                  param_grid={'C': array([1.00000000e+00, 2.78255940e+00,
7.74263683e+00, 2.15443469e+01,
5.99484250e+01, 1.66810054e+02, 4.64158883e+02, 1.29154967e+03,
3.59381366e+03, 1.00000000e+04]),
                              'penalty': ['l1', 'l2']}},
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                  scoring=None, verbose=0)
```

```
[144]: print("Pseudo R2: {:.3f}".format(grid.score(X_test, y_test)))
print('Best Penalty:', grid.best_estimator_.get_params()['penalty'])
print('Best C:', grid.best_estimator_.get_params()['C'])
```

```
test-set score: 0.986
Best Penalty: l2
Best C: 59.94842503189409
```

```
[145]: print("KFold:\n{}".format(
cross_val_score(LogisticRegression(C=59.94842503189409,penalty="l2"), X, y,
→cv=kfold).mean()))

print("StratifiedKFold:\n{}".format(
cross_val_score(LogisticRegression(C=59.94842503189409,penalty="l2"), X, y,
→cv=skfold).mean()))

print("RepeatedKFold:\n{}".format(
cross_val_score(LogisticRegression(C=59.94842503189409,penalty="l2"), X, y,
→cv=rkf).mean()))
```

```
KFold:
0.974611120980636
StratifiedKFold:
0.9879953786900199
```

RepeatedKFold:
0.9878561141706639

Logistic Regression w/ GridSearch CV w/ Scaled Data

```
[146]: grid_sc.fit(X_train_scaled, y_train)
```

```
[146]: GridSearchCV(cv=10, error_score='raise-deprecating',
                  estimator=LogisticRegression(C=1.0, class_weight=None, dual=False,
                                              fit_intercept=True,
                                              intercept_scaling=1, l1_ratio=None,
                                              max_iter=100, multi_class='warn',
                                              n_jobs=None, penalty='l2',
                                              random_state=None, solver='warn',
                                              tol=0.0001, verbose=0,
                                              warm_start=False),
                  iid='warn', n_jobs=None,
                  param_grid={'C': array([1.00000000e+00, 2.78255940e+00,
7.74263683e+00, 2.15443469e+01,
5.99484250e+01, 1.66810054e+02, 4.64158883e+02, 1.29154967e+03,
3.59381366e+03, 1.00000000e+04]),
                  'penalty': ['l1', 'l2']},
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                  scoring=None, verbose=0)
```

```
[157]: print("Pseudo R^2: {:.3f}".format(grid_sc.score(X_test_scaled, y_test)))
print('Best Penalty:', grid_sc.best_estimator_.get_params()['penalty'])
print('Best C:', grid_sc.best_estimator_.get_params()['C'])
```

Pseudo R^2: 0.995
Best Penalty: l1
Best C: 7.742636826811269

```
[158]: print("KFold:\n{}".format(
cross_val_score(LogisticRegression(C=7.742636826811269,penalty="l1"), X, y,
→cv=kfold).mean()))

print("StratifiedKFold:\n{}".format(
cross_val_score(LogisticRegression(C=7.742636826811269,penalty="l1"), X, y,
→cv=skfold).mean()))

print("RepeatedKFold:\n{}".format(
cross_val_score(LogisticRegression(C=7.742636826811269,penalty="l1"), X, y,
→cv=rkf).mean()))
```

KFold:
0.9792262687274235
StratifiedKFold:

```
0.9878401105946415
RepeatedKFold:
0.9878560904838042
```

Having performed GridSearch with respect to both penalty type and C value we were able to discover the best approach for logistic regressions with both scaled and not scaled data. The models we discovered, whose ability to predict, can be measured using the pseudo R² term predict similarly to previous models but has shown a slight improvement

0.0.18 2.4 Change the cross-validation strategy in GridSearchCV from 'stratified k-fold' to 'kfold' with shuffling. Do the parameters for models that can be tuned change? Do they change if you change the random seed of the shuffling? Or if you change the random state of the split into training and test data?

KFold w/ Shuffle False

```
[159]: grid = GridSearchCV(lgr, param_grid=hyperparameters, cv=KFold(shuffle=False,
    ↪random_state=124238920380932))
grid_sc = GridSearchCV(lgr, param_grid=hyperparameters, cv=KFold(shuffle=False,
    ↪random_state=124238920380932))

[160]: grid.fit(X_train, y_train)

[160]: GridSearchCV(cv=KFold(n_splits=3, random_state=124238920380932, shuffle=False),
    error_score='raise-deprecating',
    estimator=LogisticRegression(C=1.0, class_weight=None, dual=False,
    fit_intercept=True,
    intercept_scaling=1, l1_ratio=None,
    max_iter=100, multi_class='warn',
    n_jobs=None, penalty='l2',
    random_state=None, solver='warn',
    tol=0.0001, verbose=0,
    warm_start=False),
    iid='warn', n_jobs=None,
    param_grid={'C': array([1.00000000e+00, 2.78255940e+00,
7.74263683e+00, 2.15443469e+01,
5.99484250e+01, 1.66810054e+02, 4.64158883e+02, 1.29154967e+03,
3.59381366e+03, 1.00000000e+04])},
    'penalty': ['l1', 'l2']},
    pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
    scoring=None, verbose=0)

[161]: print("Pseudo R^2: {:.3f}".format(grid.score(X_test, y_test)))
print('Best Penalty:', grid.best_estimator_.get_params()['penalty'])
print('Best C:', grid.best_estimator_.get_params()['C'])
```

```
Pseudo R^2: 0.988
Best Penalty: l1
Best C: 166.81005372000593
```

```
[162]: grid_sc.fit(X_train_scaled, y_train)
```

```
[162]: GridSearchCV(cv=KFold(n_splits=3, random_state=124238920380932, shuffle=False),
                  error_score='raise-deprecating',
                  estimator=LogisticRegression(C=1.0, class_weight=None, dual=False,
                                                fit_intercept=True,
                                                intercept_scaling=1, l1_ratio=None,
                                                max_iter=100, multi_class='warn',
                                                n_jobs=None, penalty='l2',
                                                random_state=None, solver='warn',
                                                tol=0.0001, verbose=0,
                                                warm_start=False),
                  iid='warn', n_jobs=None,
                  param_grid={'C': array([1.00000000e+00, 2.78255940e+00,
7.74263683e+00, 2.15443469e+01,
5.99484250e+01, 1.66810054e+02, 4.64158883e+02, 1.29154967e+03,
3.59381366e+03, 1.00000000e+04]),
                              'penalty': ['l1', 'l2']}),
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                  scoring=None, verbose=0)
```

```
[163]: print("Pseudo R^2: {:.3f}".format(grid_sc.score(X_test_scaled, y_test)))
print('Best Penalty:', grid_sc.best_estimator_.get_params()['penalty'])
print('Best C:', grid_sc.best_estimator_.get_params()['C'])
```

Pseudo R²: 0.994

Best Penalty: l1

Best C: 2.7825594022071245

KFold w/ Shuffle True

```
[164]: grid = GridSearchCV(lgr, param_grid=hyperparameters, cv=KFold(shuffle=True,
↪random_state=1242389203))
grid_sc = GridSearchCV(lgr, param_grid=hyperparameters, cv=KFold(shuffle=True,
↪random_state=1242389203))
```

```
[165]: grid.fit(X_train, y_train)
```

```
[165]: GridSearchCV(cv=KFold(n_splits=3, random_state=1242389203, shuffle=True),
                  error_score='raise-deprecating',
                  estimator=LogisticRegression(C=1.0, class_weight=None, dual=False,
                                                fit_intercept=True,
                                                intercept_scaling=1, l1_ratio=None,
                                                max_iter=100, multi_class='warn',
                                                n_jobs=None, penalty='l2',
                                                random_state=None, solver='warn',
                                                tol=0.0001, verbose=0,
                                                warm_start=False),
                  iid='warn', n_jobs=None,
                  param_grid={'C': array([1.00000000e+00, 2.78255940e+00,
```



```
7.74263683e+00, 2.15443469e+01,
5.99484250e+01, 1.66810054e+02, 4.64158883e+02, 1.29154967e+03,
3.59381366e+03, 1.00000000e+04]],
    'penalty': ['l1', 'l2']],
    pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
    scoring=None, verbose=0)
```

```
[166]: print("Pseudo R^2: {:.3f}".format(grid.score(X_test, y_test)))
print('Best Penalty:', grid.best_estimator_.get_params()['penalty'])
print('Best C:', grid.best_estimator_.get_params()['C'])
```

```
Pseudo R^2: 0.987
Best Penalty: l2
Best C: 1291.5496650148827
```

```
[167]: grid_sc.fit(X_train_scaled, y_train)
```

```
[167]: GridSearchCV(cv=KFold(n_splits=3, random_state=1242389203, shuffle=True),
    error_score='raise-deprecating',
    estimator=LogisticRegression(C=1.0, class_weight=None, dual=False,
                                fit_intercept=True,
                                intercept_scaling=1, l1_ratio=None,
                                max_iter=100, multi_class='warn',
                                n_jobs=None, penalty='l2',
                                random_state=None, solver='warn',
                                tol=0.0001, verbose=0,
                                warm_start=False),
    iid='warn', n_jobs=None,
    param_grid={'C': array([1.00000000e+00, 2.78255940e+00,
7.74263683e+00, 2.15443469e+01,
5.99484250e+01, 1.66810054e+02, 4.64158883e+02, 1.29154967e+03,
3.59381366e+03, 1.00000000e+04]),
    'penalty': ['l1', 'l2']],
    pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
    scoring=None, verbose=0)
```

```
[168]: print("Pseudo R^2: {:.3f}".format(grid_sc.score(X_test_scaled, y_test)))
print('Best Penalty:', grid_sc.best_estimator_.get_params()['penalty'])
print('Best C:', grid_sc.best_estimator_.get_params()['C'])
```

```
Pseudo R^2: 0.994
Best Penalty: l2
Best C: 2.7825594022071245
```

KFold w/ New Seed

```
[169]: grid = GridSearchCV(lgr, param_grid=hyperparameters, cv=KFold(shuffle=True,
    ↪random_state=45))
```

```
grid_sc = GridSearchCV(lgr, param_grid=hyperparameters, cv=KFold(shuffle=True,
→random_state=45))
```

```
[170]: grid.fit(X_train, y_train)
```

```
[170]: GridSearchCV(cv=KFold(n_splits=3, random_state=45, shuffle=True),
          error_score='raise-deprecating',
          estimator=LogisticRegression(C=1.0, class_weight=None, dual=False,
                                         fit_intercept=True,
                                         intercept_scaling=1, l1_ratio=None,
                                         max_iter=100, multi_class='warn',
                                         n_jobs=None, penalty='l2',
                                         random_state=None, solver='warn',
                                         tol=0.0001, verbose=0,
                                         warm_start=False),
          iid='warn', n_jobs=None,
          param_grid={'C': array([1.00000000e+00, 2.78255940e+00,
7.74263683e+00, 2.15443469e+01,
5.99484250e+01, 1.66810054e+02, 4.64158883e+02, 1.29154967e+03,
3.59381366e+03, 1.00000000e+04]),
                      'penalty': ['l1', 'l2']}},
          pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
          scoring=None, verbose=0)
```

```
[171]: print("Pseudo R^2: {:.3f}".format(grid.score(X_test, y_test)))
        print('Best Penalty:', grid.best_estimator_.get_params()['penalty'])
        print('Best C:', grid.best_estimator_.get_params()['C'])
```

Pseudo R^2: 0.987

Best Penalty: l2

Best C: 464.15888336127773

```
[172]: grid_sc.fit(X_train_scaled, y_train)
```

```
[172]: GridSearchCV(cv=KFold(n_splits=3, random_state=45, shuffle=True),
          error_score='raise-deprecating',
          estimator=LogisticRegression(C=1.0, class_weight=None, dual=False,
                                         fit_intercept=True,
                                         intercept_scaling=1, l1_ratio=None,
                                         max_iter=100, multi_class='warn',
                                         n_jobs=None, penalty='l2',
                                         random_state=None, solver='warn',
                                         tol=0.0001, verbose=0,
                                         warm_start=False),
          iid='warn', n_jobs=None,
          param_grid={'C': array([1.00000000e+00, 2.78255940e+00,
7.74263683e+00, 2.15443469e+01,
5.99484250e+01, 1.66810054e+02, 4.64158883e+02, 1.29154967e+03,
3.59381366e+03, 1.00000000e+04]),
                      'penalty': ['l1', 'l2']}},
          pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
          scoring=None, verbose=0)
```

```

        'penalty': ['l1', 'l2']],
    pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
    scoring=None, verbose=0)

```

```

[173]: print("Pseudo R^2: {:.3f}".format(grid_sc.score(X_test_scaled, y_test)))
print('Best Penalty:', grid_sc.best_estimator_.get_params()['penalty'])
print('Best C:', grid_sc.best_estimator_.get_params()['C'])

```

```

Pseudo R^2: 0.995
Best Penalty: l1
Best C: 1.0

```

KFold w/ New Seed and Shuffle off

```

[174]: grid = GridSearchCV(lgr, param_grid=hyperparameters, cv=KFold(shuffle=False,
    random_state=45))
grid_sc = GridSearchCV(lgr, param_grid=hyperparameters, cv=KFold(shuffle=False,
    random_state=45))

```

```

[175]: grid.fit(X_train, y_train)

```

```

[175]: GridSearchCV(cv=KFold(n_splits=3, random_state=45, shuffle=False),
    error_score='raise-deprecating',
    estimator=LogisticRegression(C=1.0, class_weight=None, dual=False,
    fit_intercept=True,
    intercept_scaling=1, l1_ratio=None,
    max_iter=100, multi_class='warn',
    n_jobs=None, penalty='l2',
    random_state=None, solver='warn',
    tol=0.0001, verbose=0,
    warm_start=False),
    iid='warn', n_jobs=None,
    param_grid={'C': array([1.00000000e+00, 2.78255940e+00,
    7.74263683e+00, 2.15443469e+01,
    5.99484250e+01, 1.66810054e+02, 4.64158883e+02, 1.29154967e+03,
    3.59381366e+03, 1.00000000e+04]),
    'penalty': ['l1', 'l2']],
    pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
    scoring=None, verbose=0)

```

```

[473]: print("Pseudo R^2: {:.3f}".format(grid.score(X_test, y_test)))
print('Best Penalty:', grid.best_estimator_.get_params()['penalty'])
print('Best C:', grid.best_estimator_.get_params()['C'])

```

```

test-set score: 0.985
Best Penalty: l1
Best C: 21.544346900318832

```

```

[176]: grid_sc.fit(X_train_scaled, y_train)

```

```
[176]: GridSearchCV(cv=KFold(n_splits=3, random_state=45, shuffle=False),
                  error_score='raise-deprecating',
                  estimator=LogisticRegression(C=1.0, class_weight=None, dual=False,
                                              fit_intercept=True,
                                              intercept_scaling=1, l1_ratio=None,
                                              max_iter=100, multi_class='warn',
                                              n_jobs=None, penalty='l2',
                                              random_state=None, solver='warn',
                                              tol=0.0001, verbose=0,
                                              warm_start=False),
                  iid='warn', n_jobs=None,
                  param_grid={'C': array([1.00000000e+00, 2.78255940e+00,
7.74263683e+00, 2.15443469e+01,
5.99484250e+01, 1.66810054e+02, 4.64158883e+02, 1.29154967e+03,
3.59381366e+03, 1.00000000e+04]),
                  'penalty': ['l1', 'l2']},
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                  scoring=None, verbose=0)

[177]: print("Pseudo R^2: {:.3f}".format(grid_sc.score(X_test_scaled, y_test)))
print('Best Penalty:', grid_sc.best_estimator_.get_params()['penalty'])
print('Best C:', grid_sc.best_estimator_.get_params()['C'])
```

```
Pseudo R^2: 0.994
Best Penalty: l1
Best C: 2.7825594022071245
```

R²

Not Scaled, Scaled 1. KFold w/ Shuffle False -> .989,.994 2. KFold w/ Shuffle True -> .987,.994
 3. KFold w/ New Seed -> .987,.995 4. KFold w/ New Seed and Shuffle off -> .995,.994

We see that with this data there does not seem to be much change in explanatory power given the adjustments of the see and shuffling. The odd one out of the bunch appears to be 4. which is the only rendition where the not scaled data performed better than the scaled, however the difference is minimal.

0.0.19 2.5 Lastly, compare the coefficients for Logistic Regression and Penalized Logistic Regression and discuss which final model you would choose to predict new data.

```
[179]: print(lgr1.score(X_test, y_test))
print(lgr1.coef_)
```

```
0.7366153846153846
[[-0.10555434  1.26078425 -0.29515375 -4.58857718  0.71340921  1.21073104
 -2.88390237  4.91817648 -0.12439965  0.46839067  1.8779795  0.29223981]]
```

```
[180]: print(lgr2.score(X_test, y_test))
print(lgr2.coef_)
```

```
0.7366153846153846
[[-0.11178505  1.25809256 -0.29652698 -4.59653572  0.71280629  1.21350348
 -2.88391744  4.93460978 -0.13016557  0.46576363  1.88780997  0.29192203]]
```

```
[181]: print(logreg.score(X_test, y_test))
       print(logreg.coef_)
```

```
0.9876923076923076
[[ 1.41487490e+00  1.17570208e+01 -8.64211847e-01 -1.13141558e-01
  3.19822777e+01  5.31743334e-02 -6.25925503e-02 -1.99751713e+01
  8.91778445e+00  8.59114984e+00 -3.20166513e-01  2.14757181e-03]]
```

I would be inclined to use the penalized logistic regressions in this case because seldom is a model found with a R^2 so near to 1 that has no underlying issues. Multicollinearity may be at play in the non-penalized logistic regression, and this could be inspiring an artificially high R^2 value, because of this, I would inclined to reference the penalized models, as they account for multicollinearity, and have R^2 values that suggest that there may be additional features that could explain some of the variation in y , which is are more likely occurrence then having found all the independent features (x) that can explain the variation of the dependent variable (y).